



Memory, storage, and complexity (separated units)

Measured allocation (not steady-state RSS)



Versioned model material



Parameter count (count models: N/A)



Dependency and startup distinction

VOMM

Low
Python/JSON

Transformer

High
PyTorch runtime

Cold load: Transformer 3.40 ms; VOMM isolated cold load not measured