

Final next-tactic model-selection conclusion

1 · Best overall?

Transformer on accepted held-out aggregate metrics

2 · Selected for PoC?

Transformer as primary experimental predictor

3 · Why?

Material paired gains, broader metric superiority, feasible CPU runtime

4 · Main limitations?

Persistence concentration, Execution regression, weak/repeated labels

5 · Why retain VOMM?

Interpretability, Execution recall, fail-closed integration, rollback

6 · Justified claim

Corpus-level benchmark superiority

7 · Not justified

General production superiority or score fusion

Display both predictions and disagreement. Do not average scores or create tactic-dependent routing.